

<http://cs.LTH.se/EDA040>

Real-Time and Concurrent Programming

Lecture 5 (F5):

Part1: Deadlock Part 2: Messages

Klas Nilsson

2015-09-29





Content

- 1 Introduction and Definitions
 - Circular wait
- 2 Examples
 - Resource allocation graph
 - Monitor deadlock
- 3 Conditions and Analysis
 - Conditions for deadlock
 - Analysis: Resource allocation graphs
- 4 Deadlock Avoidance
- 5 Classic Example
 - The Dining Philosophers problem



Background

Mutual exclusion means that a thread can be delayed

- ▶ A thread will not be allowed to enter a critical region (using a shared resource) as long as it is occupied by another thread.
- ▶ For consistency (concurrency correctness), predictability (real-time correctness), and for efficiency (embedded computing), access of such a locked resource may not be interrupted or subject to a *roll-back*.

Hence, we have no preemption on resources (only on time as in preemptive scheduling; not to be confused).


- ▶ If the blocking never ends we have a **Deadlock** (Swe: Dödläge)

Wikipedia: Deadlock refers to a specific condition when two or more processes are each waiting for the other to release a resource, or more than two processes are waiting for resources in a circular chain.

The problem

If waiting can be or is circular:

- ▶ When several threads can be waiting for each other we have a Deadlock risk.
- ▶ When several threads are waiting for each other we have a Deadlock.



Thus, circular wait appears to be related too deadlock:

Example: deadlock with semaphores

P1

P2

```

S1.take();
S2.take();
...
S2.give();
S1.give();
  
```

```

S2.take();
S1.take();
...
S1.give();
S2.give();
  
```

P1:

S1.take();

P2:

S2.take();

S1.take();

S2.take();

blocked

blocked

- *Deadlock may occur if:* one thread performs a take, followed by a context switch (swe: trådbyte).


Example cont'd: resource allocation graph

P1

```
S1.take();  
S2.take();  
...  
S2.give();  
S1.give();
```

P2

```
S2.take();  
S1.take();  
...  
S1.give();  
S2.give();
```



Method: draw resources (boxes) and threads (circles). Draw arrows for **hold** (filled) + **wait** (outlined).


Deadlock with monitors

P1

`m1.op1();`

P2

`m2.op1();`



```

class M1 {
    synchronized void op1() {
        m2.op2();
    }
    synchronized void op2() {
        wait();
    }
}
  
```

```

class M2 {
    synchronized void op1() {
        m1.op2();
    }
    synchronized void op2() {
        wait();
    }
}
  
```


Necessary conditions for deadlock

Necessary conditions for deadlock to occur:

1. Mutual Exclusion - only one thread can access a resource at a time.
2. Hold and Wait - a thread can reserve a resource and wait for another.
3. No *resource* preemption - a thread can not be forced to release held resources.
4. Circular Wait - thread-resources dependencies must be circular.


Monitor - satisfied conditions:

1. Monitor - one thread only is allowed to enter at a time.
2. Call of an operation in a monitor from inside a monitor operation in another monitor.
3. A monitor can only be released if a thread voluntarily waits (`wait()`) or exits the monitor.
4. But 4? Must prevent circular wait that can result in deadlock.

Resource allocation graphs

Tool for detecting circular hold-wait situations and to determine under which conditions deadlock can occur.

1. Draw resources
2. Draw all hold-wait situations (arrows from each held resource to a thread marker, arrows from thread marker to resource waited for)
3. Circular? Then risk for deadlock. The number of 'hold-wait' links in the circular chain shows how many and which threads are required for deadlock.




Resource allocation graphs - example

Thread 1

```
A.take();  
B.take();  
C.take();  
C.give();  
B.give();  
A.give();
```

Thread 2

```
D.take();  
C.take();  
B.take();  
B.give();  
C.give();  
D.give();
```



Conclusion: Deadlock possible when T1 is waiting for C and T2 is simultaneously waiting for B! Circular wait!

Monitors in Concurrent Pascal

Concurrent Pascal (Per-Brinch Hansen 1979) only has monitors and the rule: *no forward references*, i.e. the program we looked at earlier is illegal:

```
class M1 {  
    synchronized void op1() {  
        // Illegal since it introduces  
        // a forward reference  
        m2.op2();  
    }  
    synchronized void op2() { ... }  
}  
  
class M2 {  
    synchronized void op1() {  
        m1.op2();  
    }  
    synchronized void op2() { ... }  
}
```

- ▶ Monitors, Semaphores, etc. are often referred to as *resources*.
- ▶ Generally: all resources is assigned a (partial) order, only allocate from lower to higher.
- ▶ M2 can call M1 but not the other way around.

Limitations in the language - a good idea?


- ▶ Deadlock impossible in Concurrent Pascal with Monitors as resources.
- ▶ Often inefficient and unpractical - might be necessary to prematurely allocate resources just to satisfy the demands on allocation order.
- ▶ Easy to implement ones own resource management using Monitors:

```

/*monitor*/ class R {
    boolean occupied;
    synchronized void request() {
        while (occupied) wait();
        occupied = true;
    }
    synchronized void release() {
        occupied = false;
        notify();
    }
}

R R1, R2;
class P1 extends Thread {
    ...
    R1.request();
    R2.request();
    ...
}
class P2 extends Thread {
    ...
    R2.request();
    R1.request();
    ...
}

```



Avoiding deadlock

- ▶ Language and library support not applicable.
- ▶ Run-time deadlock detection: Not suitable for real-time or embedded systems (useful in a generic OS, but here it would be too late).
- ▶ Instead, remove the risk for deadlock:
 - Create and analyze the resource-allocation graph.
 - Arrange the order of allocations, according to a resource ordering, preferably without extending the locking times.
 - If really necessary: add logic that prevents the dead-lock.

Notions


- ▶ Deadlock (swe: dödläge)
 - ▶ When several resources attempts to allocate the same resource one must be able to get it.
 - ▶ Bad enough if there exists an execution order such that Deadlock occurs - even if it happens only seldom. The system locks, hangs, nothing happens. Can apply to subsystems.
- ▶ Starvation (swe: svält)
 - ▶ If a thread attempts to allocate a resource it must be able to get it eventually.
 - ▶ We renounce the 'no starvation' property in favor of priority; less important activities might suffer from starvation.
- ▶ Livelock
 - ▶ Occurs when several threads attempts to allocate the same resource but none actually gets it due to the execution pattern.
 - ▶ Behaves like Deadlock, but if you study the system closely the threads actually run. They do no meaningful work though.

Dining Philosophers problem

- ▶ The life of a philosopher is boring:

```
class Philosopher
  while
    think();
    preProto();
    eat();
    postProto();
```

- ▶ Logical spaghetti:
 - ▶ Two forks are required to eat
- ▶ Solution requirements:
 - ▶ No deadlock
 - ▶ No starvation
 - ▶ Efficient



Semaphore solution 1

```

Semaphore[] fork = new MutexSem[5];
for (int i=0; i<5; i++) fork[i] = new MutexSem();

class Philosopher extends Thread {
    int i;
    Philosopher(int i) {this.i=i;}

    public void run() {
        while (true) {
            think();
            fork[i].take();
            fork[(i+1)%5].take();
            eat();
            fork[i].give();
            fork[(i+1)%5].give();
        }
    }
}

```

- ▶ Eat() acts as a critical region for a pair of forks, but for different pairs of forks
 - ▶ mutual exclusion for both
- ▶ Can this solution cause Deadlock? Two resources is required for each activity (hold-wait satisfied) so we have to make a more detailed analysis.

Philosophers 1 - deadlock?

- ▶ Draw a complete allocation graph, all Hold-Wait dependencies




- ▶ Circular → unsafe program, Deadlock can occur.
- ▶ Step 2: Can we present a scenario where deadlock occurs or prove that the situation can not occur in practice?
- ▶ Scenario: Suppose all 5 philosophers starts simultaneously, takes their left forks and then (all of them) waits for their right forks.
- ▶ This solution can thus cause deadlock.
- ▶ Can we solve the problem in a better way?
 - ▶ avoid circularity or
 - ▶ make sure that all Hold-Wait can not occur simultaneously?

Philosophers 2 - one left handed philosopher

```
Semaphore[] fork = new MutexSem[5];
for (int i=0; i<5; i++) fork[i] = new MutexSem();
```

```
class Philosopher
extends Thread {
    int i;
    Philosopher(int i)
    {this.i=i;}
    public void run() {
        while (true) {
            think();
            fork[i].take();
            fork[i+1].take();
            eat();
            fork[i].give();
            fork[i+1].give();
        }
    }
}
```

```
class LeftPhilosopher
extends Thread {
    int i;
    LeftPhilosopher(int i)
    {this.i=i;}
    public void run() {
        while (true) {
            think();
            fork[0].take();
            fork[4].take();
            eat();
            fork[4].give();
            fork[0].give();
        }
    }
}
```




- No circular dependency
- No deadlock

Philosophers 3 - only four chairs

```
Semaphore[] fork = new MutexSem[5];
for (int i=0; i<5; i++) fork[i] = new MutexSem();
Semaphore room = new CountingSem(4);
```

```
class Philosopher extends Thread {
    int i;
    Philosopher(int i) {this.i=i;}
    public void run() {
        while (true) {
            think();
            room.take();
            fork[i].take();
            fork[(i+1)%5].take();
            eat();
            fork[i].give();
            fork[(i+1)%5].give();
            room.give();
        }
    }
}
```



Complete allocation graph
cyclic, 'unsafe' as before, but:
At most four Hold-Wait can
be active simultaneously → at
least one philosopher can eat,
no deadlock possible.

Philosophers 4 - polite philosophers

A philosopher only picks up the forks and starts to eat if BOTH forks are free.

- ▶ Implemented using a monitor or a *MultistepSem*.
 - ▶ Trivially deadlock free since no Hold-Wait situations occur.
 - ▶ but, starvation possible.
- ▶ Suppose two philosophers, e.g. 1 and 3 agrees to eat alternating:
 - ▶ I.e. philosopher 1 eats until philosopher 3 has begun to eat, and the other way around
 - ▶ Now will philosopher 2 never have two forks free at the same time, i.e. **philosopher 2 will starve!!**

Part II

Message-based communication and synchronization



6 Mailboxes and messages

- Buffering and asynchronous interaction
- System aspects

7 Events and Buffers

- Messages within a program – Event objects

8 Examples

- Dataflows: Producer – Consumer



The buffering monitor as a mailbox for messages

- ▶ While monitors in general are for operations on shared data, a monitor with operations post (called by a producer thread) and fetch (called by consumer thread) comprises a data flow.
- ▶ Data can provide information and/or synchronization.
- ▶ Originally and traditionally data is then referred to as messages, and the buffer is a mailbox.
- ▶ Between threads (the same program and memory space) a message can be an Object ref.

```
class Buffer { // Providing mailbox

    synchronized void post(Object obj) {
        while (buff.size() >= maxSize) {
            wait();
        }
        if (buff.isEmpty()) notifyAll();
        buff.add(obj);
    }

    synchronized Object fetch() {
        while (buff.isEmpty()) {
            wait();
        }
        if (buff.size() >= maxSize) notifyAll();
        return buff.remove();
    }
}
```


Message sending – Mailboxes

Reasons for message-based interactions between threads:

- ▶ Producer-Consumer relations (data flows) between threads are very common
- ▶ Asymmetric synchronization (signaling) ; the producer should be allowed to continue without having to wait for the consumer.
- ▶ Transfer information – data referred to as message content.
- ▶ Thus, asynchronous communication (signaling plus data transfer) that provides **Buffering** and Thread/Activity interaction.

Additionally, for complex systems today:


- **Distribution:** Threads are, or need to be prepared for being, distributed over several computers with network communication.
- **Encapsulation:** Concurrent and real-time properties of objects (handling timeout/overrun/exceptions etc.) requires means for message passing between concurrently running objects/threads.

Situations with Producer-Consumer relations between threads are very common.


We want to achieve:

- Asymmetric synchronization - i.e. the producer should be allowed to continue without having to wait for the consumer.
- Transfer information - a message

So far we have used a Monitor/Buffer to achieve this:




We introduce a special name, Mailbox (brevlåda), for this way of using a Monitor. We draw it somewhat differently:



Systems of mailboxes

Communication between threads often forms a network of mailboxes.



- ▶ The same principle for (operating system) processes and threads.
- ▶ A thread can put a message in several mailboxes.
- ▶ A mailbox can, in Java, handle various types of messages - subclassing of message (RTEvent).
- ▶ A thread has in most cases only one mailbox it reads from (otherwise problems - fetching a message is a blocking operation).
- ▶ Message objects need to be serialized (transformed into a stream of bytes) in order to be sent to another OS process.
- ▶ Within a OS process (between threads), we can send pointers/references or a copy of the object.
- ▶ How does the receiver know that another thread does not modify the contents of a message???

IPC (Inter-Process Communication)

	Local	Distributed
'Synchronous'	Object-method call	RPC/RMI
Synchronous	Monitor-method call	Database
Asynchronous	Event buffer	Stream(pipe/file/socket)

Synchronous handling of Event

- ▶ Event model in Java (AWT, Swing Beans) basically NOT concurrent.
- ▶ Corresponding EventObject for realtime: RTEvent
- ▶ Corresponding synchronous event handling in se.lth.cs.realtime: See class documentation for RTEventListener, RTEventListenerList and JThread.
- ▶ Single-threaded 'synchronous' event handling is not a central issue in the course.

Unbounded mailbox with copy-on-send

Advantages

- ▶ Flexible code; size of buffer does not need to be decided.
- ▶ Thread safety; sent message not accessible by sender.
- ▶ The same mechanism can be used for communication between OS processes running on the same computer or different ones (distributed systems), since shared memory is not assumed.

Disadvantages

- ▶ Higher risk for running out of memory, detected later. (Memory is limited, so better set fixed bounds early.)
- ▶ Often unpractical when immediate response is required (i.e. synchronous communication).
- ▶ Increased memory use, CPU for copying, and GC work.
- ▶ Recycling via message pools difficult to implement.

We use a 'bounded buffer' in the form of `RTEventBuffer` in shared memory.

Mailbox == Monitor == Semaphore

- ▶ A Mailbox can easily be implemented using a Monitor
- ▶ Also a Semaphore is sort of a monitor.
- ▶ Suppose we only send empty messages, then a Mailbox is equivalent to a Semaphore:
 - The value of the counter of the Semaphore corresponds to the number of messages in the mailbox.
 - Send message - give()
 - Receive message - take()

All three constructions are thus equally powerful, but practical in different situations.

Events as messages

- ▶ `java.util.EventObject` comprises an event class that is suitable for messages, providing a transient (will be null outside JVM) source, referring to the sending object/thread.
- ▶ `se.lth.cs.realtime.event.RTEvent` is a subclass that, as `java.awt.InputEvent`, has a timestamp, expressing object age.
- We use such timestamped events for asynchronous communication between threads.

Note that graphics such as swing is basically single-threaded!

The RTEvent class

```
public abstract
class RTEvent extends EventObject {
    protected long timestamp;                // Creation time in ms.
    protected volatile transient Object owner; // Responsible thread.
    public RTEvent();                        // Use current Thread & TimeMillis.
    public RTEvent(Object source);           // Set source, default timestamp.
    public RTEvent(long ts);                 // Set timestamp, default source
    public RTEvent(Object source, long ts);  // Override both defaults.
    public final Object getOwner()
    public double getSeconds()
    public long getMillis()
    public long getNanos()
}
```


The RTEventBuffer class


- ▶ As for RTEvent, part of `se.lth.cs.realtime.event`

- ▶ Constructors:

```
public RTEventBuffer()  
RTEventBuffer(int maxSize)  
RTEventBuffer(int maxSize,  
               Object lock)
```

- ▶ Example with `maxSize==12` and `currSize==7`, internal attributes:

- ▶ Obtain message/event by `RTEvent fetch()` or specific final methods.
- ▶ Send message/event by `RTEvent post(RTEvent ev)` or specific final methods.



More RTEventBuffer / mailbox

Blocking and non-blocking methods for posting and fetching messages:

```
doPost(RTEvent e) // Add e to queue, blocks if the queue is full.
tryPost(RTEvent e) // Adds to the queue, without blocking if full.
doFetch()         // Fetch from queue, block if empty.
tryFetch()        // Fetch without blocking if empty.
awaitEmpty()      // Waits for buffer to become empty.
awaitFull()       // Waits for buffer to become full.
isEmpty()         // Checks if buffer is empty.
isFull()          // Checks if buffer is full.
```

The try-Post/Fetch returns null if the buffer is non-full/empty, and the supplied/next event otherwise, respectively.

The attributes are declared protected in order to make it possible to create subclasses with revised functionality.

A producer

```
class Producer extends Thread {
    Consumer receiver;
    MyMessage msg;

    public Producer(Consumer theReceiver) {
        receiver = theReceiver;
    }

    public void run() {
        while (true) {
            char c = getChar();
            msg = new MyMessage(c);
            receiver.putEvent(msg);
        }
    }
}
```

```
class MyMessage extends REvent {
    character ch;
    public MyMessage(char data) {
        super(); // Set time stamp;
        ch = data;
    }
}
```

Note: Buffering is hidden by putEvent as of the receiving thread.

A consumer

```
class Consumer extends Thread {
    RTEventBuffer mailbox;
    public Consumer(int size) {mailbox=new RTEventBuffer(size);}
    public void putEvent(RTEvent ev) {
        mailbox.post(ev); // In context of Producer
    }
    public void run() {
        RTEvent m;
        while (true) {
            m = mailbox.fetch(); // In context of Consumer
            if (m instanceof MyMessage) {
                MyMessage msg = (MyMessage) m;
                useChar(msg.ch);
            } else { ... // Handle other messages
            };
        } // ...
    }
}
```

The JThread utility class

- ▶ Part of the `se.lth.cs.realtime` package.
- ▶ Subclass of `java.lang.Thread`; thus it is a Java Thread, hence `JThread`.
- ▶ Encapsulates an `RTEventBuffer`, exposed via a public `putEvent` method.
- ▶ Default run method is a cyclic call of `perform`
- ▶ Internally the `perform` (or `run`) method refers to the `mailbox` attribute like

```
event = mailbox.doFetch();
```
- ▶ Methods `sleepUntil` and `terminate` are also provided (compare `lab1`).